Question #1 of 123

Question ID: 1536583

Your company is budgeting for the future and has requested you to create an estimate of costs related to Google Cloud resource consumption for the next quarter. What is the fastest and most efficient way for you to create this estimate?

A) Use Payment Gateway to process payments from multiple regions of Google Cloud

B) Compute costs of running a VM for a day and use that with an ML model for estimation

C) Create a data warehouse with BigQuery and store past consumption data

D) Enter the expected consumption for all resources into Google Cloud Pricing Calculator

---

Question #2 of 123

Question ID: 1536629

You are setting up your development team's Cloud infrastructure. The team is using Compute Engine virtual machines (VMs) and need to use various Google APIs and services for their workloads. The development VMs are in a single subnet and have internal IP addresses. You need to find a way to allow the VMs to access Google APIs and services. How would you do this efficiently?

A) Assign network tags for each VM on the subnet and create ingress firewall rules

B) Create static internal IP addresses for all VMs in the subnet

C) Create a customized VPC network and change its mode to auto

D) Enable Private Google Access for the subnet used by network interfaces of the VMs

---

Question #3 of 123

Question ID: 1536661

Your company performs various data processing tasks using virtual machines (VMs) deployed on its subnets using Google Cloud. You sometimes need to launch new VMs and want to be able to mitigate errors related to a used up range of IP addresses. How would you handle this error?

A) Increase the subnet's CIDR range

B) Wait for the error and redeploy failed VMs

C) Use exponential backoffs with VM restarts

D) Create several new subnets and run VMs from there

Question #4 of 123

Question ID: 1536650

You want to implement your Google Kubernetes Engine (GKE) clusters in a way that you can focus on building your apps and have Google manage the nodes and infrastructure. How will you create the appropriate management interface for this?

A) Enable the GKE API and configure an Autopilot cluster from the GKE page on the Google Cloud console

B) Enable the GKE API and then create a standard zonal cluster from the GKE page on the Google Cloud console

C) Enable the GKE API, enable gcloud CLI, and then create an alpha cluster using the `gcloud container clusters create` command

D) Enable the GKE API, enable gcloud CLI, and then create a standard regional cluster using the `gcloud container clusters create` command

---

Question #5 of 123

Question ID: 1536643

You have been working on a Compute Engine virtual machine (VM) and have performed various customizations and modifications to its persistent boot disk. You now need to quickly and efficiently create several VMs using this boot disk. What would be the fastest and most efficient way to accomplish this task?

A) Use the boot disk snapshot for creating each new VM

B) Make a custom image and then use that for creating VMs

C) Use a container image to create each VM instance

D) Create each new VM using a non-boot disk snapshot

Question #6 of 123

Question ID: 1536625

Your company does video transcoding tasks on its Google Cloud deployment. You need to upload necessary video files to a Cloud Storage bucket that is in a project you haven't created. What would be the best way to load the video data to the bucket?

A) Use Dataflow to read Pub/Sub messages and stream them to Cloud Storage

B) Create a custom Cloud Storage role and run the `storage cp` command using the gcloud CLI

C) Load CSV data from Cloud Storage and then into a table on BigQuery

D) Compress the data and export it from Cloud SQL into a MySQL instance

---

Question #7 of 123

Question ID: 1536597

You are designing a three-tier web application for a global ecommerce company using Google Cloud. You need the web application to be globally accessible with backend virtual machines (VMs) in several regions. Which of these load balancing configurations would provide the most efficient and optimal solution?

A) Use a global HTTP(S) load balancer for backend VMs and an internal HTTP(S) load balancer for receiving user requests

B) Use a global HTTP(S) load balancer for backend VMs as well as for receiving user requests

C) Use internal TCP/UDP load balancing for backend VMs and an internal HTTP(S) load balancer for receiving user requests

D) Use internal TCP/UDP load balancing for backend VMs and a global HTTP(S) load balancer for receiving user requests

---

Question #8 of 123

Question ID: 1536682

You have a service account configured to work with Cloud SQL services and you now need a user to be able to impersonate a service account so that Cloud SQL instances can be created when required. How should you engineer a solution for this?

A) Grant the principal the Cloud SQL Admin role

B) Grant the principal the IAM Editor role

C) Grant the principal the Compute Admin role

D) Grant the principal the Service Account User role

Question #9 of 123

Question ID: 1536657

You work for a company that provides database solutions for various financial services companies using Cloud SQL on Google Cloud. You need to perform an inherently risky operation on the database immediately and you cannot wait for the automated backup window to pass. What action will you need to perform before the unsafe operation?

A) Configure a cross-region read replica for the Cloud SQL database

B) Migrate the database to a Firestore document database

C) Perform an on-demand backup for the Cloud SQL database

D) Perform a point-in-time recovery and create a new instance

---

Question #10 of 123

Question ID: 1536613

You are running a Google Kubernetes Engine (GKE) cluster, which is native to your VPC and uses internal IP addresses. You want to allow specific nodes inside of the cluster to send outbound traffic to the internet. How would you implement this?

A) Create a managed instance group for the VMs in the cluster

B) Use Cloud NAT for outbound internet connectivity for nodes

C) Create external IP addresses for each node in the cluster

D) Change the cluster type from zonal to regional

Question #11 of 123

Question ID: 1536569

You work for a company that is moving its operations onto Google Cloud. The company has multiple departments and projects in them. You need to map the existing organizational structure correctly into a Google Cloud hierarchy. When building a resource hierarchy on Google Cloud, what is the recommended top-down order for the hierarchy?

A) Organization, Folders, Projects, Resources

B) Resources, Folders, Projects, Organization

C) Resources, Projects, Folders, Organization

D) Organization, Projects, Folders, Resources

Question #12 of 123

Question ID: 1536596

You need to select a low latency load balancing option for your Google Cloud system that requires user data to be stored in a single geographical area for regulatory compliance with traffic entering the load balancer from the internet through a single public IP address while ensuring high levels of performance. Which load balancing option should you use in this scenario?

A) Internal HTTP(S) load balancer

B) Internal TCP/UDP load balancing

C) Global external HTTP(S) load balancer

D) Regional external HTTP(S) load balancer

Question #13 of 123

Question ID: 1536628

The company has several different projects and needs to find a way for resources within the projects to communicate with each other using internal IP addresses. You also need to ensure that this system can be managed centrally in an efficient manner. How would you implement this?

A) Create a custom VPC network and change its mode to auto

B) Create a global subnet and add all the cloud resources to it

C) Move all the resources across projects into a new project in a new organization

D) Designate a host project and specify eligible resources in service projects

Question #14 of 123

Question ID: 1536630

Your company uses Compute Engine resources on Google Cloud. It needs to enforce access control for its virtual machines (VMs) that are being used for developing web applications. You need to allow incoming TCP traffic to specific VMs in the company's VPC network. How will you meet this requirement?

A) Create a new egress firewall rule with the default target and create a new IP subnet within the VPC

B) Add a target tag of http-server to specific VMs and set an ingress firewall rule for it

C) Use the default target on specified VMs and set an egress firewall rule for it

D) Create a new service account to access data on the VPC using APIs

Question #15 of 123

Question ID: 1582590

You are running a web service using the Google Cloud platform and utilizing APIs for performing various business tasks. You notice that API requests are failing with the error code `413 REQUEST ENTITY TOO LARGE`. How would you handle this issue?

A) Request a higher quota limit using the Google Cloud console

B) Deploy more Compute Engine VM instances

C) Create an autoscaled managed instance group using an instance template

D) Free up resources from your deployment and restart all applications

Question #16 of 123

Question ID: 1536617

Your workloads comprise several containerized applications that you need to run using GKE. You want to ensure that the containerized application is deployed using Google's recommended best practices including collecting GKE cluster credentials and addition of recommended labels to the resource file. How would you do this?

A) Use Cloud Run's gke-deploy for deploying the application to a GKE cluster

B) Create external IP addresses for each node in the GKE cluster

C) Create outbound internet connections to Google servers using Cloud NAT

D) Use kubectl builder to deploy the application to a GKE cluster

---

Question #17 of 123

Question ID: 1536666

Your Google Cloud deployment has several running projects, each of which have resources generating log entries. To better understand the system and find ways to improve its performance, you need to set up a system where specific logs are routed automatically to BigQuery and the rest are discarded. How will you do this?

A) Create an exclusion filter for specific logs for routing to a BigQuery dataset

B) Create an inclusion filter for specific logs for routing to a BigQuery dataset

C) Remove all exclusion filters for the sink and allow automatic routing

D) Find routed log entries from Cloud Logging in the BigQuery console

---

Question #18 of 123

Question ID: 1536646

You have several Kubernetes clusters running on your Google Cloud deployment and you need to analyze the resources and metrics related to the clusters, including all nodes, pods, and services. How will you accomplish this task?

A) Tag a Docker image with a repository name and push it to Artifact Registry

B) Use GKE with the Autopilot mode enabled for all existing clusters

C) Access the GKE dashboard from the Monitoring section of the Google Cloud console

D) Create a managed instance group for the VMs in the GKE cluster

---

Question #19 of 123

Question ID: 1536627

Your company has a financial accounting application that is used by several consulting firms. The application is exposed to the firms through an IP address, which provides access to the server that runs the application. You need to find a cost-effective solution to ensure that there is no downtime to the system and IP address changes on the DNS caused by server changes are automatically managed. How would you engineer this solution?

A) Use Cloud DNS for assigning a static external IP address for the application

B) Distribute application data using a CDN with Cloud Storage for data storage

C) Create a custom script that updates the domain's IP address during server crashes

D) Use an IP address of your choice using the Bring Your Own IP (BYOIP) system

---

Question #20 of 123

Question ID: 1536649

You are using the Google Cloud platform to deploy an application in Google Kubernetes Engine (GKE). Your application has several replicas, each of which has an application container. How will you ensure that the number of replicas are adjusted automatically based on average CPU utilization?

A) Create a VerticalPodAutoscaler object and set the updateMode to Auto

B) Use both the Horizontal Pod Autoscaler along with the Vertical Pod Autoscaler for CPU metrics

C) Change the labels of all Pods currently running on the GKE cluster

D) Create a HorizontalPodAutoscaler object using the `kubectl autoscale` command

---

Question #21 of 123

Question ID: 1536631

You work for a global digital services organization that uses Google Cloud Compute Engine virtual machines (VMs) for its workloads. You need these VMs to connect with a partner company's VMs that use Amazon Web Services (AWS) Elastic Compute Cloud (EC2) instances. This connection needs to be made with high levels of availability and security over the internet. How would you implement this solution?

A) Create a global subnet and add both cloud service provider resources to it

B) Create a customized VPC network on Google Cloud and change its mode to auto

C) Connect to AWS using an IPSec VPN connection

D) Create an ingress rule for the Google Cloud VMs with AWS as the target tag

---

Question #22 of 123

Question ID: 1536568

Your company is running multiple projects that utilize Google Cloud resources and needs to manage access to them through policies. You want to ensure that the policies specified for the organization are enforced on all the resources running under specific projects in addition to the policies that are placed for the resources themselves. How would you engineer this solution?

A) Create a resource hierarchy with folders, projects, and resources in several organizations

B) Create one all-encompassing policy and apply it to all the necessary resources

C) Create a resource hierarchy with folders, projects, and resources in one organization

D) Create an individual policy for each resource that has the necessary policy conditions

---

Question #23 of 123

Question ID: 1536642

Your company uses Compute Engine virtual machines (VMs) on Google Cloud for machine learning (ML) applications, and you need to upgrade the processing power of the VMs. You have successfully created a new VM using a public image and have added GPUs to it. What action should you take to allow applications to use the GPUs correctly?

A) Perform a hard reset of the VM before use

B) Use the `regions describe` command to verify regional GPU quota

C) Install the NVIDIA CUDA Toolkit on the VM

D) Retrieve the instance ID of the VM and rename it

Question #24 of 123

Question ID: 1536611

You are using Google Kubernetes Engine (GKE) clusters in your Google Cloud deployment. You have installed kubectl and you need to interact with and configure the GKE clusters. What will you need to do for this?

A) Create Tag keys and values and attach the Tags to a cluster

B) Use a custom client for interacting with GKE clusters

C) Install the gke-gcloud-auth-plugin to communicate with GKE clusters

D) Create key-value pairs for goog-gke- GKE cluster labels

Question #25 of 123

Question ID: 1536600

You are designing the Google Cloud system for a web application development company. You need to implement a frontend for a wide number of globally distributed backend instances to ensure resiliency of the application. How would you implement this as efficiently as possible?

A) Provide direct connectivity to the internal load balancer for each global user

B) Connect the global load balancer directly to backend VM instances

C) Use a single anycast IP address for access to the application

D) Use a separate IP address for each backend instance in the application

---

Question #26 of 123

Question ID: 1536689

You are supervising the operations on a Google Cloud deployment and need to ensure that resources on the cloud deployment have the correct metadata as per company policy. You need to know which users ran APIs that changed the metadata of various Google Cloud resources. How will you find this information?

A) Set up ingestion and storage for Policy Denied audit logs in Cloud Logging

B) Configure Admin Activity audit logs for metadata changes in resources

C) Record changes to cloud resource configurations using System Event audit logs

D) Enable and check Data Access audit logs using Logs Explorer

---

Question #27 of 123

Question ID: 1536672

You are running Google Kubernetes Engine (GKE) clusters for various processing tasks in a Google Cloud deployment and your team has a new member who needs to perform cluster updates using gcloud CLI or the Google Cloud console. How will you ensure the team member can perform this job?

A) Instruct the project owner to grant the Kubernetes Engine Cluster Admin role to the new team member

B) Instruct the project owner to grant the Kubernetes Engine Viewer role to the new team member

C) Instruct the project owner to grant the Compute Viewer role to the new team member

D) Instruct the project owner to grant the basic IAM viewer role to the new team member

---

Question #28 of 123

Question ID: 1536684

You are administering the Google Cloud deployment for your company. You need to engineer a system that allows a specific user the privileges to handle emergencies in the production system by delegating troubleshooting tasks. How will you make this happen?

A) Grant the user the Basic IAM Editor role

B) Grant the user the Service Account User role

C) Grant the user the Basic IAM Owner role

D) Grant the user the Service Account Token Creator role

Question #29 of 123

Question ID: 1536636

A software development company has various departments working on its products, which include teams for developers, testers, and analytics experts. You need to create a solution to deploy each team's environment on Google Cloud and maintain stability through Google's recommended best practices. What approach would you take for this?

A) Create a single template for deployment with Cloud Foundation Toolkit and use Terraform for deployment

B) Create a separate bucket for each team on Cloud Storage with custom IAM permissions

C) Deploy each environment using gcloud commands inside of a shell script

D) Create a unique Terraform configuration for each environment and update these regularly

---

Question #30 of 123

Question ID: 1536681

You are in the process of creating service accounts for your applications to use that are running on virtual machines (VMs) on your Google Cloud deployment. Which of the following will you need to do to use service accounts efficiently?

A) Utilize the IAM API for auditing service accounts, keys, and allow policies

B) Use the default service accounts whenever possible

C) Keep all team members on the service account's allow policy

D) Maintain public accessibility for all existing external keys

---

Question #31 of 123

Question ID: 1536615

Your company performs big data processing and is looking to use Google Cloud for managing its clusters and nodes. You need an open-source system to automatically manage, deploy, and scale container-based applications with granular control over OS and memory configurations. What would you do for this requirement?

A) Create Tag keys and values and attach the Tags to all clusters

B) Create applications using Firestore and integrate them with BigQuery

C) Utilize GKE in Standard mode for the applications

D) Configure and deploy custom machine type VM instances

Question #32 of 123

Question ID: 1536603

Your company uses Google Cloud Compute Engine to run its web applications. You have created several Linux virtual machine (VM) instances to process its workloads. What will you need to do to configure the applications running on these instances?

A) Configure and deploy new custom machine type VM instances

B) Use in-memory caching for the applications with Cloud Spanner for data storage

C) Use Remote Desktop Protocol (RDP) to connect to the VMs over the internet

D) Connect to the instances using the Google Cloud CLI and Secure Shell (SSH)

Question #33 of 123

Question ID: 1536675

You are creating a system for users on your company's Google Cloud to access and use various resources like Compute Engine virtual machine (VM) instances and Cloud Storage buckets. You need to grant permissions to a principal, but existing roles have other permissions as well that the principal must not receive for security reasons. How would you engineer a solution for this scenario?

A) Create a custom role for the principal

B) Find a predefined role that performs the necessary tasks

C) Grant the IAM Owner role to the principal

D) Assign a Basic IAM role to the principal

Question #34 of 123

Question ID: 1536593

Your company has large volumes of customer data that it needs to store on Google Cloud Storage. You want to find the most affordable storage solution for durable storage data, which may be accessed less than once a year, for legal and compliance requirements. What storage class on Cloud Storage would you use for this?

A) Nearline

B) Standard

C) Coldline

D) Archive

Question #35 of 123

Question ID: 1582586

A software development company has asked you to create a Google Cloud deployment that can aid its engineering teams in meeting business goals. You need to set up a system to ensure Site Reliability Engineering (SRE) best practices and to ensure you have low latency in your applications. How would you accomplish this?

A) Deploy containerized applications in a preferred language using Cloud Run

B) Build an enterprise level data warehousing system with BigQuery

C) Use Cloud Trace, Debugger, and Profiler to identify code issues in applications

D) Use VMware Engine to move applications onto Google Cloud

Question #36 of 123

Question ID: 1536644

You have several virtual machines (VMs) inside of a managed instance group for processing data processing tasks. You need to ensure optimum usage of Google Cloud resources on your cloud through varying VM processing levels. How would you specify an autoscaling parameter to scale VM instances for optimizing usage costs?

A) Create a new instance template with the machine type based on average CPU usage

B) Create a new unmanaged instance group and add all VMs to it

C) Change the machine type for VMs in the instance group after disabling autoscaling

D) Specify the target metric for average CPU utilization for all VMs in the group

Question #37 of 123

Question ID: 1536659

Your company is creating a data warehousing system of petabytes of data that can be accessed by data scientists globally for their work. You need to design a solution that allows the scientists to access the data using SQL at any time with the total usage of up to 40 minutes a day. What would be the most efficient and cost-effective way to deliver this solution?

A) Store the data using Cloud Storage standard class and use a script for queries

B) Use BigQuery for both storage and querying of the data

C) Use Cloud BigTable for both storage and querying of the data

D) Create a Kubernetes cluster using GKE and use containers for processing data

---

Question #38 of 123

Question ID: 1582583

You work for a software development company that uses the Google Cloud platform. Presently, the company has several thousand employee accounts on Google Cloud. You need to find a way to create a user group that contains all the technical writers in the company. What would be the fastest and most efficient way to do this?

A) Create a dynamic group using a membership query for technical writers

B) Send out an email to all employees asking them to add themselves to a group

C) Use Google's automatic identity federation feature with Active Directory

D) Search employee roles for technical writers and add them to a Google group

---

Question #39 of 123

Question ID: 1536654

Your company performs digital video editing and transcoding using Google Cloud Storage for storing raw videos. You want to delete all videos that were uploaded to the bucket before January 1, 2023. What would be the most efficient way to do this?

A) Change the storage class for all stored objects on the bucket to Coldline

B) Perform data replication using Cloud Backup and Disaster Recovery (DR)

C) Create a lifecycle configuration for the storage bucket from the Cloud Storage Buckets page on the Google Cloud console

D) Use the Google Cloud console to find all objects uploaded before January 1, 2023, and delete them

---

Question #40 of 123

Question ID: 1536606

You need to transfer files to your existing virtual machine (VM) using Cloud Storage. For this, you have created a Cloud Storage bucket and uploaded your files there. Now, you need to connect to the VM using a third-party app and download the files from the bucket. What will you need to do for this?

A) Use ssh-keygen and generate a new SSH key pair and connect to the VM

B) Create policies for Object Lifecycle Management on the Cloud Storage bucket

C) Create an autoscaled managed VM group using an instance template

D) Create a script using Cloud Function triggers based on Cloud Storage events

---

Question #41 of 123

Question ID: 1582582

A company is performing a Google Cloud adoption for its network security application. It requires various security officers to view the application's network traffic data and use BigQuery to run reports. You must implement this requirement in Google Cloud in the most efficient way possible. How would you engineer this solution?

A) Create a custom view-only role for each security officer for viewing BigQuery results

B) Assign the **roles/bigquery.dataviewer** role to each security officer

C) Create a group specifically for the security officers and assign them the **roles/bigquery.dataviewer** role

D) Create a group specifically for the security officers and assign them the **roles/viewer** role

---

Question #42 of 123

Question ID: 1582587

You have a batch data processing application that can withstand occasional instance stoppage. You need to find the most cost-optimized Compute Engine system for this application that does not add an additional workload on your existing virtual machine (VM) instances. How would you set up this solution?

A) Configure and deploy custom machine type VM instances

B) Configure a lifecycle policy for data objects used by the application

C) Use and run preemptible VMs

D) Use Cloud Trace for finding latency within the application

---

Question #43 of 123

Question ID: 1536612

You work for a web services company that has large web applications running on Google Cloud. You are currently using zonal clusters and notice downtime in the system each time you perform IP address rotation. You need to maintain high availability in the system. How would you do this?

A) Use a custom client for interacting with the clusters

B) Create new regional clusters for the deployment

C) Create several new zonal clusters for the system

D) Change the cluster type from zonal to regional

Question #44 of 123

Question ID: 1536663

You work as a Cloud Engineer for a web services team using Google Cloud and you are currently running Apache Web Server on a Compute Engine virtual machine (VM) instance. You need to be notified when the traffic rate on the web server goes beyond 10 KiB/s so you can manage the server accordingly. How would you do this?

A) Create a custom metric that measures the number of processes running on the server

B) Create an alerting policy and set the Threshold value for the alert trigger to 10000

C) Create a log-based alert that notifies you when the web server reboots

D) Create a new network tag for the VM instance and create an egress firewall rule for it

---

Question #45 of 123

Question ID: 1536653

Your company uses Google Cloud Storage for storing all its organizational data. You are asked to provide a solution to ensure that the data stored on Cloud Storage buckets is secured from threats. Which of these should you recommend as a way to secure objects in Cloud Storage buckets at the enterprise level?

A) Create a lifecycle configuration for the storage bucket from the Cloud Storage Buckets page on the Google Cloud console

B) Enable uniform bucket-level access and domain-restricted sharing

C) Replicate all stored data using Cloud Backup and Disaster Recovery (DR)

D) Provide authorization to Google Cloud to secure all the company's data stored on Cloud Storage

---

Question #46 of 123

Question ID: 1536605

Your company provides a highly available online banking service to clients globally. You need to ensure cost-optimized operation on your multi-zone Google Cloud deployment by having virtual machine (VM) instances deleted when their CPU use falls below an average value. You also need to be able to repeat this configuration for many VMs in the future. What would be the most efficient way to engineer this solution?

A) Configure each VM instance and write a custom script for analyzing utilization

B) Use preemptible VMs within a managed instance group

C) Create a zonal managed instance group

D) Create an autoscaled managed instance group using an instance template

---

Question #47 of 123

Question ID: 1536581

Your company uses Google Cloud for several of its Internet of Things (IoT) application development projects. The finance department has asked you to find a way to view detailed data related to resource usage, cost estimates, and pricing through an entire workday and be able to view this data visually. What should you do first to meet this requirement?

A) Use Cloud Profiler for profiling consumption of resources across your production applications

B) Utilize Cloud Monitoring for gathering relevant metadata, events, and metrics from components of applications

C) Create a project, enable billing on it, and set permissions for the project and Cloud Billing account

D) Set up Cloud Logging to store, search and analyze app log data, and set up alerts and notifications

---

Question #48 of 123

Question ID: 1536620

Your company is creating a new social media and microblogging app for users on the Google Cloud platform. Users can upload image and video content to their accounts using the app. You need to detect images with inappropriate content and automatically blur them as they are uploaded to a bucket on Cloud Storage. How will you engineer this solution?

A) Use ImageMagick to manually blur each inappropriate image using IAM user roles

B) Create a data processing service with Cloud Run and invoke it using Pub/Sub

C) Have an authorized user inspect uploaded images and upload appropriate ones to a new bucket on Cloud Storage

D) Change the storage class of the bucket on Cloud Storage to Coldline

---

Question #49 of 123

Question ID: 1537015

You are the cloud administrator for a company that has workloads that are beyond a typical user's lifecycle. These workloads utilize Google Cloud APIs. You want a user to fully manage the process of authentication and authorization for data access using APIs for non-human users on your Google Cloud deployment. What will you need to do for this?

A) Grant the roles/resourcemanager.folderAdmin role to the user

B) Grant the roles/resourcemanager.projectIamAdmin role to the user

C) Grant the roles/editor role to the user

D) Grant the roles/iam.serviceAccountAdmin role to the user

Question #50 of 123

Question ID: 1536618

You have deployed several containerized applications to your GKE clusters. You need to ensure the clusters are operating as expected. For this, you need to view key metrics related to the cluster including open issues and CPU and memory utilization from a single dashboard. How would you do this?

A) Create Tag keys and values and attach the tags to clusters

B) Run GKE for all clusters with the Autopilot mode enabled

C) Deploy a monitoring application for each cluster using Cloud Run's gke-deploy

D) Set integration settings for Cloud Monitoring and Cloud Logging for the cluster from the console

---

Question #51 of 123

Question ID: 1536652

Your company has a large on-premises data center as well as a Google Cloud deployment. Both setups utilize containerized applications for various client services. You need a way to deploy your workloads across multiple clouds and hybrid environments while maintaining a consistency of usage experience. How should you engineer this solution?

A) Use Cloud Run with Knative API deploying the workloads on GCP

B) Integrate legacy systems into a cloud-based build system using Cloud Build

C) Migrate existing VMs to new containers using VMware Engine

D) Deploy the workloads using Cloud Run for Anthos

Question #52 of 123

Question ID: 1536623

You work for a mobile app development company that is building a new chat app using Google Cloud. This will require you to create a system for storing a user's profile and chat rooms. How would you engineer this solution?

A) Use Cloud SQL to create a fully managed relational database to store app data

B) Create a Cloud Firestore database and store data in documents and collections

C) Create a data warehouse using BigQuery and use SQL for data querying

D) Implement an OLTP database for online transactions using Cloud Spanner

Question #53 of 123

Question ID: 1582593

You are running several Dataproc jobs on your Google Cloud deployment and you need to look into the details of the jobs to troubleshoot issues related to job scheduling delays. How would you investigate these issues?

A) Submit a new BigQuery job using existing Dataproc job IDs and view job history from the BigQuery page on the Google Cloud console

B) Run the `gcloud dataflow jobs list` command with the relevant creation dates and region

C) Change the built-in default logs and metrics collection configuration for Ops Agent

D) Run the `gcloud dataproc jobs describe` gcloud command with the relevant region and job IDs

Question #54 of 123

Question ID: 1582594

There are various logs stored on your Google Cloud deployment. You need to reduce the costs incurred by storing logs that are generated by various services in order to optimize the operational cost of your Google Cloud deployment. How can you accomplish this?

A) Enable Data Access audit logs for Cloud Storage

B) Disable System Event audit logs from being written

C) Create an exclusion filter for Policy Denied violations

D) Create a script for excluding Admin Activity audit logs

Question #55 of 123

Question ID: 1536619

You are using Cloud Run to deliver a web-based service to clients. Your service needs to minimize latency and reduce the number of cold starts. What would you do to ensure that your system stays in a “warm” state and keeps some container instances available even during zero traffic?

A) Enable idle instances for the service

B) Manually scale the instance count to match usage spikes

C) Perform traffic splitting among multiple revisions of the function

D) Modify the latest revision of the service with optimized code

---

Question #56 of 123

Question ID: 1536690

There have been reports that resources on the Google Cloud deployment have incurred cost overrun issues. You need to see if there have been new virtual machines (VMs) created recently and identify the user account that created them. How will you accomplish this?

A) List the total policy violations made by users through the Policy Denied audit logs

B) Check Data Access audit logs using Logs Explorer for API calls to read metadata

C) Generate System Event audit logs for tracking resource configuration changes

D) Create a query in Logs Explorer to search on Admin Activity audit logs

---

Question #57 of 123

Question ID: 1536635

Your company needs to maintain a strong compliance and security posture while deploying its new Google Cloud environment. How will you provision a full production environment and ensure that key governance and security controls are implemented instantly in your new cloud environment?

A) Create a foundational system using CFT templates and use IaC

B) Specify a Custom role for each employee in the organization

C) Create a resource hierarchy with folders, projects, and resources in one organization

D) Create organizational policies for all the project's resources

---

Question #58 of 123

Question ID: 1582595

Your company needs to maintain its adherence to various federal regulations. Many of these regulations pertain to storage of personally identifiable information (PII) and you need to create a compliance report that contains complete details of data access and accountability on your Google Cloud deployment. How will you perform this task?

A) View job history from the BigQuery page on the Google Cloud console

B) Access Cloud Audit logs using Logs Explorer

C) Configure the built-in default logs and metrics collection configuration for Ops Agent

D) Create a policy for Object Lifecycle Management on all Cloud Storage buckets

---

Question #59 of 123

Question ID: 1536648

You are designing a new multi-tier web app with a frontend and backend on Google Cloud. The web app has a frontend application that you are running using Kubernetes. What is the best way to increase the web app's availability?

A) Deploy the application as StatefulObjects and use PersistentVolume resources with dynamic provisioning

B) Deploy multiple replicas of the application into uniform Pods using the Deployment controller

C) Deploy the application as StatefulObjects and integrate Filestore as a NFS persistent volume for storage

D) Deploy a single instance of the application into multiple unique Pods using the Deployment controller

---

Question #60 of 123

Question ID: 1536578

Your company is taking on several software development projects for various vendors. The company's teams use Google Cloud projects with various compute and storage resources for application development. You need to find a way for the company to manage the expenses related to all projects efficiently. What would you need to do for this?

A) Create labels for all Google Cloud project resources based on vendors

B) Create organizational policies for all Google Cloud project resources

C) Create a separate Cloud Billing account for each Google Cloud project

D) Create a single Cloud Billing account for all Google Cloud projects

---

Question #61 of 123

Question ID: 1536599

Your company needs to implement a Google Cloud deployment to perform data processing services for its global clients. How would you design the layout of locations for the company's cloud resources on the Google Cloud network to ensure high availability within a region?

A) Keep all web tiers and backend instances in the us-east1-b zone

B) Keep a single web tier and use database tiers without backend VMs

C) Keep one backend instance in the us-east1-a zone and another in the us-east1-b zone

D) Keep each backend instance in the us-east1-a zone

Question #62 of 123

Question ID: 1536667

You are running a web application on Google Cloud and you need to know the time taken by the application to handle application and user requests. You also need to know the time taken to complete RPC calls when requests are handled by the application. How will you accomplish this?

A) Create an inclusion filter for the application's logs and route them to BigQuery

B) Inspect the Service Health page for specific incidents related to Cloud Trace

C) View trace data for the application from the Cloud Trace Overview page

D) Check the most current status for Cloud Trace from the Google Cloud Status page

---

Question #63 of 123

Question ID: 1536601

Your company is a global telecom firm using Google Cloud. You need to configure a solution for serving a Domain Name Server (DNS) zone that can be viewed on the internet. You need to ensure that the zone has high availability and low latency. How would you engineer this?

A) Create a DNS database using BigQuery

B) Use Internal DNS for Compute Engine VMs and load balancers

C) Use Cloud DNS as a server for public zones on the internet

D) Run domain controllers using the Managed Service for Microsoft Active Directory

---

Question #64 of 123

Question ID: 1536688

You are looking into security incidents related to misallocation of user privileges on your Google Cloud deployment and you need to understand who made the changes to the access privileges and when this was done. How would you do this?

A) Search Admin Activity audit logs for actions made to change IAM permissions

B) List the total policy violations made by users through the Policy Denied audit logs

C) Generate System Event audit logs for tracking resource configuration changes

D) Find a list of all user-run API calls to modify resources using Data Access audit logs

---

Question #65 of 123

Question ID: 1536579

Your company is currently handling several projects for various global clients, each of which will be billed in a different currency. You have created a new Google Cloud project for each client's project. How would you manage expenditures for a specific project on Google Cloud?

A) Link each created project to a Cloud Billing account

B) Use multi-region buckets for data storage across all projects

C) Set resource quotas and perform throttling across projects

D) Create organizational policies for all the project's resources

---

Question #66 of 123

Question ID: 1536686

The Google Cloud administrator has assigned several tickets to you, each of which are issues related to services not being available to users. To resolve these issues, you need to find details of all the situations where specific resource access on Google Cloud was not permitted due to a user not having sufficient privileges. How will you do this?

A) Check user-run API calls in Data Access audit logs

B) Generate System Event audit logs

C) View Policy Denied audit logs

D) Configure Admin Activity audit logs

Question #67 of 123

Question ID: 1536674

You need to manage access and permissions for your company's Google Cloud object storage and database management. For this, you need to grant a role to a user that enables viewing Cloud Storage buckets and administering databases on Cloud SQL. How will you choose the most appropriate predefined role for this?

A) Locate a single predefined role that enables permissions for both the services

B) Grant the roles/editor role to the principal for both Cloud SQL and Cloud Storage

C) Find two suitable predefined roles, one with permissions for Cloud SQL and the other for Cloud Storage

D) Grant a basic IAM Admin role for Cloud SQL databases and use that for Cloud Storage as well

---

Question #68 of 123

Question ID: 1536590

Your company provides ecommerce solutions for several global clients and is expanding its services using Google Cloud. You need to create a NoSQL serverless document database for real-time inventories and user profiles that enable clients to watch, document, and receive real-time updates on it. How would you engineer this solution?

A) Create a data preparation system using Dataprep

B) Create application logic using SQL and store data using BigQuery

C) Build a scalable, fully managed service for Redis using Memorystore

D) Create the application using Firestore

Question #69 of 123

Question ID: 1536655

Your company is using a Google Cloud Storage bucket for storing all its log data. To save on storage costs, you are asked to configure the system so that all objects that are more than a year old are moved to the Coldline storage class. You need this settings to be applied to both current objects as well as future objects stored on the bucket. How would you do this?

A) Create a new Cloud Storage bucket with the Coldline storage class and move objects there that are a year old

B) Set the storage class for the bucket to Coldline

C) Enable uniform bucket-level access and domain-restricted sharing

D) Set a lifecycle configuration for the Cloud Storage bucket for all objects

---

Question #70 of 123

Question ID: 1536656

Your company performs data processing tasks for various US banks. The financial department needs to know the expected cost of using Google Cloud Storage for archiving data in a specific US region. What would be the best way to create an accurate estimate for this?

A) Compute the cost of storing data in buckets for a day and use that with an ML model for estimation

B) Use the Pricing Calculator with the storage class set for Coldline

C) Set the storage class for Cloud Storage buckets to Coldline and view the Cloud Billing Reports page

D) Create a lifecycle configuration for all storage buckets on Cloud Storage to move all objects to Coldline

---

Question #71 of 123

Question ID: 1536651

Your Cloud Run service is working with a database system that has a limit on how many concurrent connections it can handle. Which of these scaling parameters will you need to configure for autoscaling container instances using Cloud Run?

A) The maximum container instances setting

B) The number of rollbacks to function revisions

C) The maximum concurrent requests per instance

D) The average CPU utilization of scheduled instances

Question #72 of 123

Question ID: 1536668

You are asked to handle the billing needs of all running projects on your company's Google Cloud deployment. You need to ensure that there is no chance of resources being consumed beyond an acceptable amount due to either a cyberattack or misuse. How can you ensure this?

A) Perform resource labeling and analyze billing data

B) Perform encryption of data at rest across all projects

C) Specify quotas for all project resources

D) Create budgets and configure alerts for all projects

Question #73 of 123

Question ID: 1536664

Your development team used Compute Engine for its coding work and is currently developing a new transcoding application. You need to retrieve data on the number of background processes spawned by the application running on a virtual machine (VM) instance. How would you do this?

A) Locate a built-in metric within Cloud Monitoring to measure the number of background processes

B) Use OpenCensus to create a custom metric to measure the number of background processes

C) View trace data for the application from the Cloud Trace Overview page

D) Migrate the existing VM to a new container using VMware Engine

Question #74 of 123

Question ID: 1536670

You are working as a Google Cloud engineer with a company providing Software as a Service (SaaS) systems with multi-tenant clusters. You need to ensure that your nodes are safe from cyberattacks. How would you ensure this?

A) Configure the Cloud Identity account to use SSO

B) Use gVisor with GKE Sandbox

C) Create a MongoDB connection

D) Set encryption of data at rest across all projects

Question #75 of 123

Question ID: 1536632

Your data processing company uses Google Cloud for its workloads that work with confidential data. You need to balance encrypted traffic from the internet securely onto virtual machines (VMs) on your VPC network on Google Cloud. Which load balancing system would be the optimal choice for this scenario?

A) Internal TCP/UDP Load Balancing

B) External TCP Proxy Load Balancing

C) External SSL Proxy Load Balancing

D) Internal HTTP(S) Load Balancing

Question #76 of 123

Question ID: 1536622

Your company is designing a new photo sharing app that will allow its users to upload pictures to the app and share them with other users. The app is designed to ensure cost-optimized storage with uploaded pictures first converted to a binary format. What would be the most efficient and cost-effective way to do this?

A) Use Visual Inspection API for finding and localizing issues inside the uploaded images

B) Use Cloud Storage for image uploads and use Cloud Functions to run binary conversion code each time there is an upload

C) Use Memorystore for image storage and use Cloud Functions to run binary conversion code each time there is an upload

D) Use Cloud Vision for automatic analysis and classification of uploaded images

---

Question #77 of 123

Question ID: 1537014

Your company is looking to store all its multimedia data on Google Cloud platform. You need a system that keeps this data available for access once a quarter in the most economical way possible. What solution would you choose for this requirement?

A) Create a Cloud Storage bucket and set the default storage class to Coldline

B) Use Dataflow to process batches of data and perform analytics using BigQuery

C) Create a Cloud Function that accesses NoSQL data on Cloud BigTable

D) Perform ETL operations on the data using Dataproc and use Kerberos for security

---

Question #78 of 123

Question ID: 1536683

You are administering a project on Google Cloud and need to perform tasks to resolve various issues in the cloud environment and you need to obtain temporary elevated privileges to perform debugging activities on various Compute Engine instances. How would you obtain elevated privileges that you need temporarily for this scenario?

A) Assign the compute.instances.stop permission to the user

B) Grant the Compute Admin role to the user

C) Grant the principal the IAM Editor role

D) Create an OAuth2.0 access token for the related service account

---

Question #79 of 123

Question ID: 1536595

You work for a digital publishing company that has an on-premises data center with storage resources for audio and video. You need to find a cost optimized and low maintenance solution to connect virtual machine (VM) instances on a virtual private cloud (VPC) to the on-premises resources. How would you engineer this solution?

A) Connect the on-premises resources to the VPC using a VPN

B) Build a direct connection to each VM using software routing on Compute Engine VMs

C) Provide public IP addresses to each VM for access through the internet

D) Connect the on-premises resources to the VPC using Cloud Interconnect

---

Question #80 of 123

Question ID: 1536639

You are creating a new virtual machine (VM) on a subnet and you want to ensure it cannot be accessed from outside the network. How will you accomplish this?

A) Edit the VM configuration and change its machine type for scale-out optimized workloads

B) Add a network interface to the VM without an external IP address

C) Assign a network tag for the VM and create an egress firewall rule for it

D) Enable guest attributes on the VM for storing host keys

Question #81 of 123

Question ID: 1536665

Your Google Cloud system has several Compute Engine instances that generate logs containing key system information and metrics. You want these metrics to be analyzed to further optimize the system and you need to find a way to route logs generated through Google Cloud to third-party systems for analysis. How would you engineer this solution?

A) Configure the sink destination as a Pub/Sub topic from the Logs Router page on the Google Cloud console

B) Configure the sink destination as a Cloud Storage bucket from the Logs Router page on the Google Cloud console

C) Modify the built-in default logs and metrics collection configuration for Ops Agent

D) Create a custom metric that measures the number of logs generated by the Compute Engine VMs

---

Question #82 of 123

Question ID: 1536647

You are using Kubernetes clusters on your Google Cloud platform to provide a Service to clients over the internet and you need to increase the processing power of a node pool on your GKE cluster. What would be the most appropriate way to do this?

A) Change the labels of the Pods running on the cluster

B) Upgrade the machine type for specific nodes of the cluster individually

C) Configure the entire node pool for an upgraded machine type

D) Deploy the Service to a specific node pool

Question #83 of 123

Question ID: 1536592

You are a cloud engineer for a database and data warehousing services company that has its operations running on Google Cloud. You need to engineer a block storage solution for its advanced database workloads requiring random access and throughput in bulk. Which storage option would best serve this requirement?

A) Balanced persistent disks

B) Local SSDs

C) Standard persistent disks

D) Extreme persistent disks

Question #84 of 123

Question ID: 1582584

Your company is a data analytics firm that has a large number of employees. The workflows require employees to use several specific third-party apps on the cloud. You need to find a way to allow new users to access the specified apps and have the apps and user profiles automatically deprovisioned when a user leaves the company. What would be the most efficient way to implement this solution?

A) Enable access to third-party apps for all users and delete unused apps from the cloud

B) Create a group for each app in use and add only relevant users to that group

C) Create a dynamic group using a membership query for specific apps

D) Activate SAML for SSO to provision users and use Cloud Identity to sync users to apps

---

Question #85 of 123

Question ID: 1536598

You are creating a computing architecture for a software development firm specializing in web applications. You need to perform traffic routing to internal virtual machines (VMs) that are inside the same regional virtual private cloud (VPC) network. Which load balancing option would be most suitable for systems within a VPC?

A) Regional external HTTP(S) load balancer

B) External TCP/UDP network load balancing

C) Global external HTTP(S) load balancer

D) Internal TCP/UDP load balancing

Question #86 of 123

Question ID: 1536587

You are setting up a cloud computing solution for your company's online shopping application and you need to find a way to distribute user transactions across multiple virtual machines (VMs) whenever the application usage increases. You need this to be performed automatically and in real time. How would you achieve this solution in the most cost-optimized manner?

A) Use Cloud SQL with a memory cache and perform polling on VMs for unprocessed transactions

B) Route transactions to BigQuery and perform polling on VMs to identify unprocessed transactions

C) Keep transactions on Cloud BigTable and perform polling on VMs for processed transactions

D) Route transactions to Pub/Sub and use VMs in a managed instance group for processing

---

Question #87 of 123

Question ID: 1536634

Your company is developing a new blogging app using WordPress and has asked you to find a suitable pre-built solution for it. You have found several packages that can be used on Cloud Marketplace. Before you deploy a package, you need to ensure that the product meets your specific requirements. How will you do this?

A) Have your development team develop a similar package and deploy both to compare

B) Deploy the package first and then perform an update on it using Cloud Marketplace

C) Check details of the VM types, OS, and cost estimates from the product details page

D) Change the number of virtual machine (VM) instances configured with the package

---

Question #88 of 123

Question ID: 1536662

You work as a Google Cloud consultant for a graphics development company. The company has created image editing applications that users can access over the internet. These applications are running on Compute Engine virtual machine (VM) instances. As per Google Cloud best practices, what is the best way to expose the applications to users?

A) Increase the subnet's CIDR range to cover all VMs

B) Use domain names that are mapped to IP addresses using Cloud DNS

C) Set an internal DNS name for each VM instance and share these with users

D) Allow Compute Engine to set DNS names to VMs and share these with users

---

Question #89 of 123

Question ID: 1536588

You need to implement a cloud storage solution for your company on Google Cloud. All data needs to be stored using standard storage except for data older than one year which needs to be archived automatically for regulatory compliance. Also, data older than five years must be deleted automatically. What is the most efficient and cost-effective way to architect this solution?

A) Create a script to run every day that deletes data that is five years old and moves data that is one year old to a separate bucket for archiving

B) Create a separate bucket for deleted objects and run a script to move data that is older than five years to that bucket

C) Create a script that checks objects daily and marks data that is one year old as “To Archive” and deletes data that is five years old

D) Create policies for Object Lifecycle management

Question #90 of 123

Question ID: 1536673

You are creating an architecture describing the structure of resource access and permissions your company requires and you need to correctly identify the relevant principals for creating the access management system. Which of these would you specify as principals?

A) View actions on buckets

B) Service accounts

C) Storage resources

D) Allow policies

Question #91 of 123

Question ID: 1582591

Your company is using microservices based on Docker containers for adding, updating, and deleting user data for its global video streaming application. The number of operations can vary from a few to many millions. You want to deploy this application using Google Cloud in a way that minimizes overhead and has massive scalability. How would you do this?

A) Set up BigQuery to store data from the Docker containers

B) Use GKE for deploying the Docker containers

C) Use an instance group for launching every Docker container

D) Use Cloud Run for deploying the Docker containers

Question #92 of 123

Question ID: 1536633

Your company wants to accelerate its digital transformation using Google Cloud. You have workloads that need virtual machines (VMs) and cloud-based storage, but your team lacks the expertise for using Compute Engine and Cloud Storage. You need a pre-built solution that is based on your resource consumption needs and not manually configured VM instances or storage buckets. What would be your plan of action for this scenario?

A) Use and run preemptible VMs for your workloads with Object Lifecycle Management

B) Enter your expected consumption for all resources into Google Cloud Pricing Calculator

C) Access the Cloud Marketplace page from the Google Cloud console and select a suitable software package

D) Use Cloud Profiler for profiling consumption of resources across your production applications

---

Question #93 of 123

Question ID: 1582592

You are running a Compute Engine VM instance in your VPC network for ad hoc batch processing tasks and you are looking for ways to optimize the cost of your cloud resources. What should you do if you want to keep the VM instance and not be charged for it?

A) Run the `gcloud compute instances suspend` command for the instance

B) Delete the instance and create a new one from the Images page on the Google Cloud console

C) Stop the VM and change its machine type from the VM instances page on the Google Cloud console

D) Run the `gcloud compute instances create` command and set enable-guest-attributes=TRUE for enabling guest attributes on the VM

Question #94 of 123

Question ID: 1582581

You work for an Internet of Things (IoT) application development firm that is using Google Cloud's Pub/Sub service. You have specified the allow policies in such a way that `bruce@example.com` has the role of Editor for the Content_Dev project, and Amy has the role of Publisher for topic_1 in the same project. You have now set an allow policy on topic_1, which grants Bruce the Viewer role for the topic. How would this affect the way policies stand for this project?

A) Bruce has the Viewer role for topic_1 and Amy has the Publisher role for topic_1

B) Amy has the Editor role for topic_1 and Bruce has the Viewer role for topic_1

C) Both Bruce and Amy have the Viewer role for topic_1

D) Bruce has the Editor role for topic_1 and Amy has the Publisher role for topic_1

Question #95 of 123

Question ID: 1536676

You are in the process of delegating responsibilities for running tasks on a company's Google Cloud deployment and you need to allow a user to manage access to all the Google Cloud resources in the system. Which IAM predefined role will you need to grant the user for this?

A) roles/iam.securityAdmin

B) roles/resourcemanager.organizationAdmin

C) roles/resourcemanager.folderAdmin

D) roles/resourcemanager.projectIamAdmin

Question #96 of 123

Question ID: 1536626

Your company is undergoing a cloud adoption process and moving its entire operations from an on-premises system to the Google Cloud platform. It currently has 300TB of data that needs to be put on its new cloud system, and data uploads using the internet fall within the company's data usage plans. What would be the most cost-optimized way of performing this data transfer?

A) Upload the data to Cloud Storage using broadband

B) Send the data to Google using a Transfer Appliance

C) Use Cloud Backup and DR to replicate the data

D) Perform data transfer using a Dedicated Interconnect link

Question #97 of 123

Question ID: 1536607

You have a data processing application on Google Cloud that utilizes hundreds of virtual machines (VMs). You want to get relevant system metrics from existing and new VMs associated with your Google Cloud project. You also need the installation and maintenance of the agent to be performed across all VMs in the most efficient way possible. What would you do for this requirement?

A) Create an Agent Policy using gcloud CLI

B) Install Ops Agent on each individual VM and upgrade them manually

C) Create an autoscaled managed VM group using an instance template

D) Create a script using Cloud Function triggers based on VM instance events

---

Question #98 of 123

Question ID: 1536582

You work for a web applications development company that is using Google Cloud for its operations. The development team has recently created a new application using AppEngine and wants to deploy it using automation. What would you need to do to provide this functionality?

A) Install Java and then download and execute the Google Cloud CLI installer

B) Install Ruby and then create labels for all Google Cloud project resources

C) Install C++ and then use multi-region buckets for data storage across all projects

D) Install Python and then download and execute the Google Cloud CLI installer

---

Question #99 of 123

Question ID: 1582589

Your Google Cloud deployment is running various virtual machine (VM) instances using Compute Engine. These instances need to be monitored closely to ensure optimal system health and performance. For this, you have an Ops Agent running on each instance, but you need to customize specific groups of metrics that should be collected by the Ops Agent. How will you achieve this requirement?

A) Enable authorization for the VM instances for running Ops Agent

B) Modify the built-in default logs and metrics collection configuration for Ops Agent

C) Install the gcloud CLI beta component and create an Agent Policy for the VM instances

D) Create a file of overrides to change the Ops Agent collection configuration for metrics

---

Question #100 of 123

Question ID: 1536602

You are creating a website for a geoinformatics company using Google Cloud. You have a new domain name registered and have created a Compute Engine Linux virtual machine (VM) instance. You have also installed Apache web server on the instance. How should you configure the system so that the domain name will be hosted using the VM instance?

A) Set the Zone type to Private on the Create a DNS Zone page

B) Update the name servers for the domain in Google Domains with NS records

C) Create a new VM instance for processing domain name requests

D) Enter the VM instance's external IP address on the web server's default page

---

Question #101 of 123

Question ID: 1536594

You work for a digital publishing company that is looking to build its data storage solutions on Google Cloud. You need to implement a system to store short video clips that are accessed frequently through the month and long videos that are accessed once a quarter at most. What would be the most cost-optimized solution for this requirement?

A) Store the video clips using the Standard object storage class and longer videos with the Nearline class

B) Store the video clips using the Nearline object storage class and longer videos with the Standard class

C) Store the video clips using the Coldline object storage class and longer videos with the Nearline class

D) Store the video clips using the Archive object storage class and longer videos with the Standard class

---

Question #102 of 123

Question ID: 1582580

You are asked to apply organizational policies to your company's Google Cloud resource hierarchy on Cloud Storage. You need to ensure that Amy, who is in charge of data processing, should be able to view and delete any uploaded files but not delete any buckets. Additionally, you need to allow users to upload files to buckets without being able to view or delete other users' uploads. How would you implement this solution?

A) Add Amy to the user group `file_upload_team@example.com` and grant the group the role of Storage Object Creator

B) Grant the role of Storage Object Creator to Amy and Storage Object Admin to the user group `file_upload_team@example.com`

C) Grant the role of Storage Object Creator to both Amy and the user group `file_upload_team@example.com`

D) Grant the role of Storage Object Admin to Amy and Storage Object Creator to the user group `file_upload_team@example.com`

Question #103 of 123

Question ID: 1536616

You work for a data processing company that uses Kubernetes clusters using Google Kubernetes Engine (GKE). The system is designed so that a backend Service Pods need to communicate with frontend Pods without any interruptions. How can you ensure this?

A) Perform grouping of backend Pods through a Service that accepts communication from frontend Pods

B) Use a static IP address with a DNS setting that the frontend Pods can use

C) Allow frontend Pods to contact backend Pods using static internal IP addresses

D) Ensure all Pods communicate using a static external IP address

---

Question #104 of 123

Question ID: 1536641

Your company is a network application development firm and is using Compute Engine on Google Cloud to create a new file sharing app. You need to ensure that connections made to virtual machines (VMs) are protected against vulnerabilities. Which of these is a best practice for securely connecting to VM instances on your Google Cloud VPC?

A) Change the machine type of the VMs for network-centric workloads

B) Allow VMs to be accessed on the public internet

C) Enable guest attributes on the VMs for storing host keys

D) Perform a hard reset of the VMs for memory cleaning

Question #105 of 123

Question ID: 1536604

You need to configure your Cloud deployment so that the system functionality can be ensured when any restarts are performed on VMs for maintenance. How would you configure the maintenance behavior of VM instances on your Cloud?

A) Assign new persistent disks to each existing VM instance

B) Create fresh VM instances from a boot disk image

C) Specify different Google Cloud regions for VM instances

D) Set availability policies using the metadata of the VM instances

---

Question #106 of 123

Question ID: 1536680

You work for a digital record label that uses Google Cloud Storage buckets to store user uploaded audio files and you are setting up an access management system for the music file processing application. You need to use the principle of least privilege to allow the application's service account to check the format of the audio files stored on the Cloud Storage bucket. How would you engineer this solution?

A) Grant the service account the Storage Object Viewer role for the Cloud Storage bucket

B) Grant the service account the Storage Object Admin role for the Cloud Storage bucket

C) Grant the service account the Storage Object Creator role for the Cloud Storage bucket

D) Grant the service account the Storage Insights Viewer role for the Cloud Storage bucket

---

Question #107 of 123

Question ID: 1536580

You work for a graphics development company that uses Google Cloud for its graphics projects. You need to ensure that a programmatic notification is sent to the key stakeholders on Slack and project billing is disabled when resource usage costs for the project go beyond an expected limit. How would you implement this solution?

A) Utilize Pub/Sub for connecting a topic to the budget

B) Perform analytics on usage data stored in a BigQuery data warehouse

C) Customize email notifications using Cloud Monitoring

D) Create labels for all Google Cloud project resources based on costs

---

Question #108 of 123

Question ID: 1536660

You run a Google Cloud deployment for data processing using Compute Engine virtual machines (VMs). You have several VMs inside of a subnet and you need to ensure that you keep the same IP address for a specific VM resource even if the resource is deleted and remade. How would you do this?

A) Assign a network tag for the VM and create an ingress firewall rule for it

B) Enable Private Google Access for the subnet used by network interfaces of the VM

C) Reserve a static internal IP address from the configured IP range for the subnet

D) Add a new dual stack subnet inside the VPC network and add the resource to it

---

Question #109 of 123

Question ID: 1536637

Your company has various development team members who work with multiple Google Cloud resources. You want to use Kubernetes API for managing your Google Cloud resources to reduce load on your development team. How would you do this?

A) Install Config Connector using its GKE add-on

B) Set Cloud NAT for outbound internet connectivity for nodes

C) Create Tag keys and values and attach the tags to a cluster

D) Create key-value pairs for goog-gke- GKE cluster labels

Question #110 of 123

Question ID: 1536591

Your company performs data processing services using Google Cloud. You are working with a business-critical data processing application that needs very high availability. You need to ensure that any zonal failure is mitigated instantly by having storage attached to a standby virtual machine (VM). What block storage option would you use in this scenario?

A) Cloud BigTable

B) Regional balanced persistent disks

C) Local SSDs

D) Zonal persistent disks

Question #111 of 123

Question ID: 1536678

You have several applications running on a Compute Engine Virtual Machine (VM) instance that need to access Google Cloud resources, such as Cloud Storage buckets and BigQuery tables. For this, you need to engineer a system where applications can be authenticated for accessing resources and successfully making authorized API calls. What would you need to do for this?

A) Grant a basic IAM role of Editor to the VM users

B) Assign the compute.instances.stop permission to VM users

C) Use role recommendations to find insights on existing principals

D) Create and attach a service account to the VM

Question #112 of 123

Question ID: 1536610

You have a running project on Google Cloud. Your business operations are mission critical and cannot tolerate any application downtime. You need to find a way to get notified immediately when your project is near a quota limit. How would you implement this?

A) Create a custom alert for project resources using Cloud Monitoring API

B) Free up resources from your deployment and restart all applications

C) Configure each VM instance and write a custom script for generating quota alerts

D) Change built-in default metrics collection configurations for Ops Agent

---

Question #113 of 123

Question ID: 1536584

Your company uses Cloud Storage and Filestore for its data storage. You need to find a way to monitor these systems and ensure that any action made to change the data is in line with existing regulatory compliance. You also need to ensure that this system is implemented optimally and performs as close to real time as possible. What should you do for this requirement?

A) Run a security script using a triggering system in every data storage system

B) Run a security script after creating a script in Python for collecting and analyzing data storage logs

C) Run a security script after sending each data update query first to an App Engine app

D) Run a security script using Cloud Function triggers based on Cloud Function events

---

Question #114 of 123

Question ID: 1536677

There have been concerns raised to the company's management that various users have accessed resources that they should not be allowed to. You need to strengthen the security configuration of your company's Google Cloud deployment with regard to access permissions for all of its users. Which of these strategies would you use to enforce this?

A) Give basic IAM Owner roles to key principals in the organization

B) Review and apply role recommendations for existing principals

C) Allow the iam.roles.create role for all existing cloud users

D) Grant predefined Viewer roles to all Google Cloud users for all resources

---

Question #115 of 123

Question ID: 1582585

Your company has a new project with an application that needs machine learning (ML)-based analytics services along with compute and storage resources. You decide to use the necessary services from Google Cloud. What should you do to allow your development team to access these services over the public internet?

A) Enable Cloud APIs for the project using the Google Cloud console

B) Create a group for each service to be used and add only relevant users to that group

C) Use Google's automatic identity federation feature with Active Directory

D) Activate billing for the project using Google Cloud console

---

Question #116 of 123

Question ID: 1536671

You are setting permissions for various computing and storage resources on your company's Google Cloud deployment, and you need to know how a particular resource inherits access permissions based on the resource hierarchy in the organization. How will you find this information?

A) Use gcloud CLI to get the effective IAM policy for the resource

B) Enable automatic user provisioning for creating, modifying, and deleting user identities

C) Create a resource hierarchy with folders, projects, and resources in one organization

D) Use an external identity provider for workforce authorization

Question #117 of 123

Question ID: 1536669

You work for a large multinational corporation that has offices around the world all using Google Cloud for various computational tasks. You are asked to gather and analyze all the charges of cloud resource consumption for purposes of cost optimization. What would be the most efficient way to do this as per Google's best practices?

A) Use resource labels, export billing data to BigQuery, and use Looker Studio for visual analysis

B) Store text files with resource data on Cloud Storage and run a script for data analysis

C) Write a customized script using C++ for analyzing cloud resources and query the results using Cloud SQL

D) Use tags for resources, export billing data to BigQuery, and use Data Studio for analysis

Question #118 of 123

Question ID: 1536573

Your company is doing a project on its Google Cloud deployment for the US Department of Defense. The team working on the project has members from different departments in the company. You need to ensure that only certain team members have access to sensitive information in the project by specifying the necessary Identity and Access Management (IAM) roles. What would be the most optimal way to implement this solution?

A) Create groups, specify an IAM Predefined role to the groups, and add users to them

B) Create groups, set a basic role for each group, and then add users to each group

C) Create groups, specify Custom roles for each group, and then add users to the groups

D) Specify a Custom role for each team member

Question #119 of 123

Question ID: 1536589

Your company has designed a new application that performs body temperature and pulse monitoring for global users through various Internet of Things (IoT) devices like smartwatches. It needs a storage system that can store up to 20 petabytes (PBs) of data through a simple schema that allows fast reads and writes of small data chunks. What would be the most economical and efficient solution for this requirement?

A) Use Cloud BigTable for data storage and create the application logic in any programming language

B) Use in-memory caching with Cloud Spanner for data storage

C) Perform data distribution using a CDN while using Cloud Storage for data storage

D) Create application logic using SQL and store data using BigQuery

---

Question #120 of 123

Question ID: 1536638

Your company performs data analytics work using a virtual machine (VM) on its VPC on Google Cloud. You need to change the configuration of the VM by increasing its vCPU and memory and add GPUs to meet changing workloads. How would you do this?

A) Suspend the VM by shutting down the guest OS on it

B) Delete the VM instance and create a new one using a boot disk image

C) Reset the VM to clean its memory state

D) Stop the VM and make changes to its machine type

Question #121 of 123

Question ID: 1536614

Your company uses Kubernetes clusters for various applications on Google Cloud. It needs you to architect a solution to manage the complete infrastructure of a cluster with no need for monitoring or configuration to ensure cost optimization. How would you engineer this solution?

A) Keep a data table containing node data on BigQuery and use a script for resource management

B) Use GKE with the Autopilot mode

C) Use a managed instance group for Compute Engine with auto scaling enabled

D) Perform manual configuration of cluster capacity and avail discounts

---

Question #122 of 123

Question ID: 1536640

Your development team has been working on a particular Compute Engine virtual machine (VM) instance for application development work. You need to purge the memory of the VM and bring it to its initial state. What is the recommended approach for this?

A) Enable guest attributes on the VMs for storing host keys

B) Perform a reset of the VM instance

C) Delete the VM instance and create a new one using a boot disk image

D) Suspend the VM by shutting down the guest OS on it

Question #123 of 123

Question ID: 1582588

Your company uses Google Cloud for advanced scientific computing workloads. You need to deploy virtual machines (VMs) that are fully optimized for your company's workloads with the exact amount of memory and CPU that is required by scientific research driven applications. What will you need to do to configure this requirement in the most optimal way?

A) Create VMs using a custom machine type

B) Create VMs using a predefined VM type as it is

C) Create a preemptible VM type

D) Enable auto start and stop with predefined VM types